



UNIVERSITY OF
ARKANSAS

LSTM Recurrent Neural Networks

Jiahui Chen

University of Arkansas

References: Wei MSU, colah github repo

Introduction to RNN

Recurrent neural networks (RNN) were developed by John Hopfield (1982). In 1993, an RNN used 1000 layers in time. It is a very active research topic and one of the top ten methods in machine learning.

Motivation:


Temporal patterns or correlations in sequential data, such as speech recognition, handwriting recognition, credit card crime prevention, drug optimization, subject of a researcher's next paper etc., can be used to make better predictions.

Introduction to RNN


Sequence modeling of sequential data:
Design objective

- Deal with **variable-length** sequences
- Track **long-term** dependencies
- Maintain information about the **order**
- **Share parameters** across the sequence

A memory cell of regular RNN



Simple RNN – Elman Network




$$h_t = \sigma_h(C_h X_t + U h_{t-1} + b_h)$$

Weight: C_h, U


$$\hat{y}_t = \sigma_y(C_y h_t + b_y)$$

Hidden: h_t

Simple RNN – Elman Network



Simple RNN – Jordan Network




Loss Function of RNN

- $h_t = \sigma_h(C_h X_t + U h_{t-1} + b_h)$ tanh is used
- $\hat{y}_t = \sigma_y(C_y h_t + b_y)$ softmax is used
- Cross entropy loss function for each step
- $l_t(\hat{y}_t, y_t) = -\sum_i^M y_t^{(i)} \log \hat{y}_t^{(i)}$
- Accumulated total error
- $L(\hat{y}_t, y_t) = \sum_t l_t(\hat{y}_t, y_t)$

Backpropagation

- Chain rule: $\frac{\partial L}{\partial C_{ij}} = \sum_t \frac{\partial l_t(\hat{y}_t, y_t)}{\partial C_{ij}}$
- $\frac{\partial l_t(\hat{y}_t, y_t)}{\partial C_{ij}} = \sum_{k=0}^t \frac{\partial l_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial C_{ij}}$



Vanishing Gradient Issue

- $\frac{\partial l_t(\hat{y}_t, y_t)}{\partial C_{ij}} = \sum_{k=0}^t \frac{\partial l_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial h_t} \prod_{m=k+1}^t \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial C_{ij}}$
- **Gradients from far away steps vanish.** As a result, RNN does not maintain the long-range memory, which is needed in many situations.

Vanishing Gradient Issue

Solutions


- Use appropriate treatment of weight matrices: initialization and normalization
- Use ReLu, which is linear for positive values
- Use Long-Short-Term Memory
- Use Gated Recurrent Units
- Use Residual network (ResNets)
- Use Deep belief network, etc.

Vanishing/Exploding Gradient Problem

- Backpropagated errors multiply at each layer, resulting in exponential decay (if derivative is small) or growth (if derivative is large).
- Makes it very difficult train deep networks, or simple recurrent networks over many time steps.

Long Distance Dependencies


- It is very difficult to train RNNs to retain information over many time steps
- This make is very difficult to learn RNNs that handle long-distance dependencies, such as subject-verb agreement.



Long Short Term Memory


- LSTM networks, add additional gating units in each memory cell.
 - Forget gate
 - Input gate
 - Output gate
- Prevents vanishing/exploding gradient problem and allows network to retain state information over longer periods of time.

LSTM Network Architecture



Cell State

- Maintains a vector C_t that is the same dimensionality as the hidden state, h_t
- Information can be added or deleted from this state vector via the forget and input gates.




Cell State Example

- Want to remember person & number of a subject noun so that it can be checked to agree with the person & number of verb when it is eventually encountered.
- Forget gate will remove existing information of a prior subject when a new one is encountered.
- Input gate "adds" in the information for the new subject.

Forget Gate


- Forget gate computes a 0-1 value using a logistic sigmoid output function from the input, x_t , and the current hidden state, h_t :
- Multiplicatively combined with cell state, "forgetting" information where the gate outputs something close to 0.



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$


Hyperbolic Tangent Units

- Tanh can be used as an alternative nonlinear function to the sigmoid logistic (0-1) output function.
- Used to produce thresholded output between -1 and 1 .



Input Gate

- First, determine which entries in the cell state to update by computing 0-1 sigmoid output.
- Then determine what amount to add/subtract from these entries by computing a tanh output (valued -1 to 1) function of the input and hidden state.

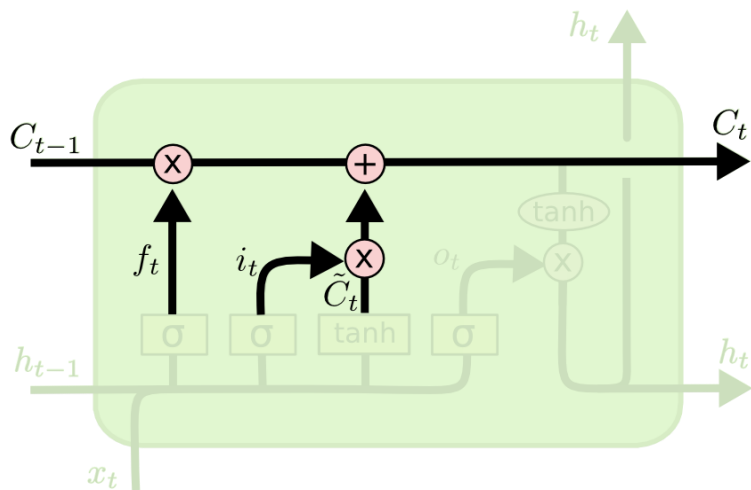


$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Updating the Cell State


Cell state is updated by using component-wise vector multiply to "forget" and vector addition to "input" new information.



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Output Gate

- Hidden state is updated based on a "filtered" version of the cell state, scaled to -1 to 1 using \tanh .
- Output gate computes a sigmoid function of the input and current hidden state to determine which elements of the cell state to "output".

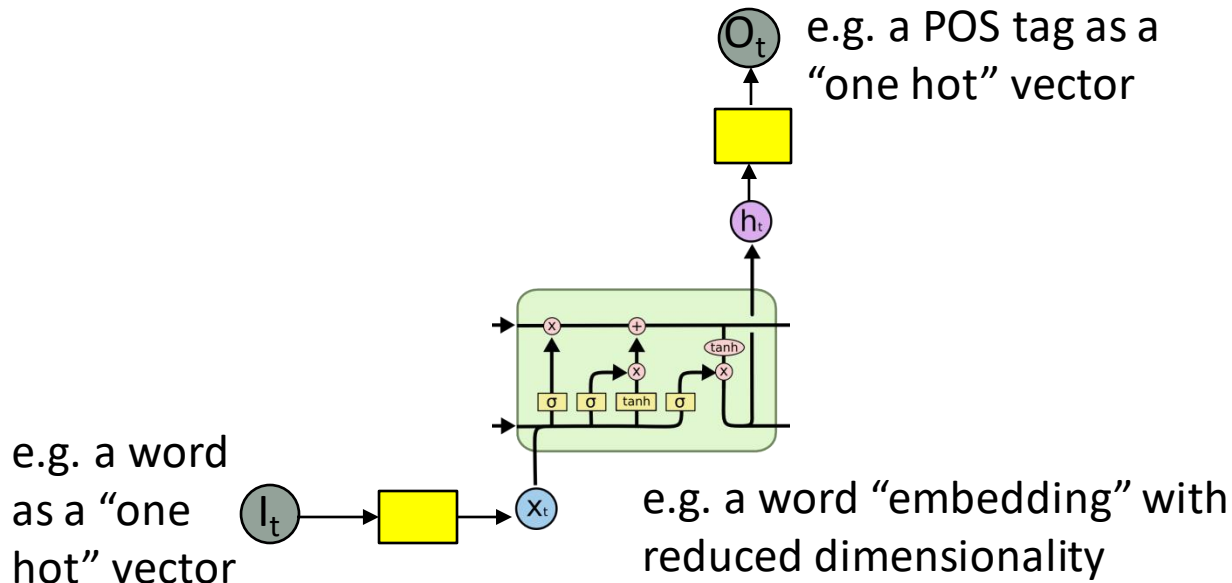


$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

Overall Network Architecture

Single or multilayer networks can compute LSTM inputs from problem inputs and problem outputs from LSTM outputs.



LSTM Training


- Trainable with backprop derivatives such as:
 - Stochastic gradient descent (randomize order of examples in each epoch) with momentum (bias weight changes to continue in same direction as last update).
 - ADAM optimizer (Kingma & Ma, 2015)
- Each cell has many parameters (W_f W_i W_o W_c)
 - Generally requires lots of training data.
 - Requires lots of compute time that exploits GPU clusters.

General Problems Solved with LSTMs

- Sequence labeling
 - Train with supervised output at each time step computed using a single or multilayer network that maps the hidden state (h_t) to an output vector (O_t).
- Language modeling
 - Train to predict next input ($O_t = I_{t+1}$)
- Sequence (e.g. text) classification
 - Train a single or multilayer network that maps the final hidden state (h_n) to an output vector (O).

Sequence to Sequence Transduction

- Encoder/Decoder framework maps one sequence to a "deep vector" then another LSTM maps this vector to an output sequence.




- Train model "end to end" on I/O pairs of sequences.

General Problems Solved with LSTMs

- Sequence labeling
 - Train with supervised output at each time step computed using a single or multilayer network that maps the hidden state (h_t) to an output vector (O_t).
- Language modeling
 - Train to predict next input ($O_t = I_{t+1}$)
- Sequence (e.g. text) classification
 - Train a single or multilayer network that maps the final hidden state (h_n) to an output vector (O).

Sequence to Sequence Transduction

- Encoder/Decoder framework maps one sequence to a "deep vector" then another LSTM maps this vector to an output sequence.



- Train model "end to end" on I/O pairs of sequences.

LSTM Application Architectures

one to many






Image Captioning

many to one




Video Activity Recog
Text Classification

many to many



Video Captioning
Machine Translation

many to many

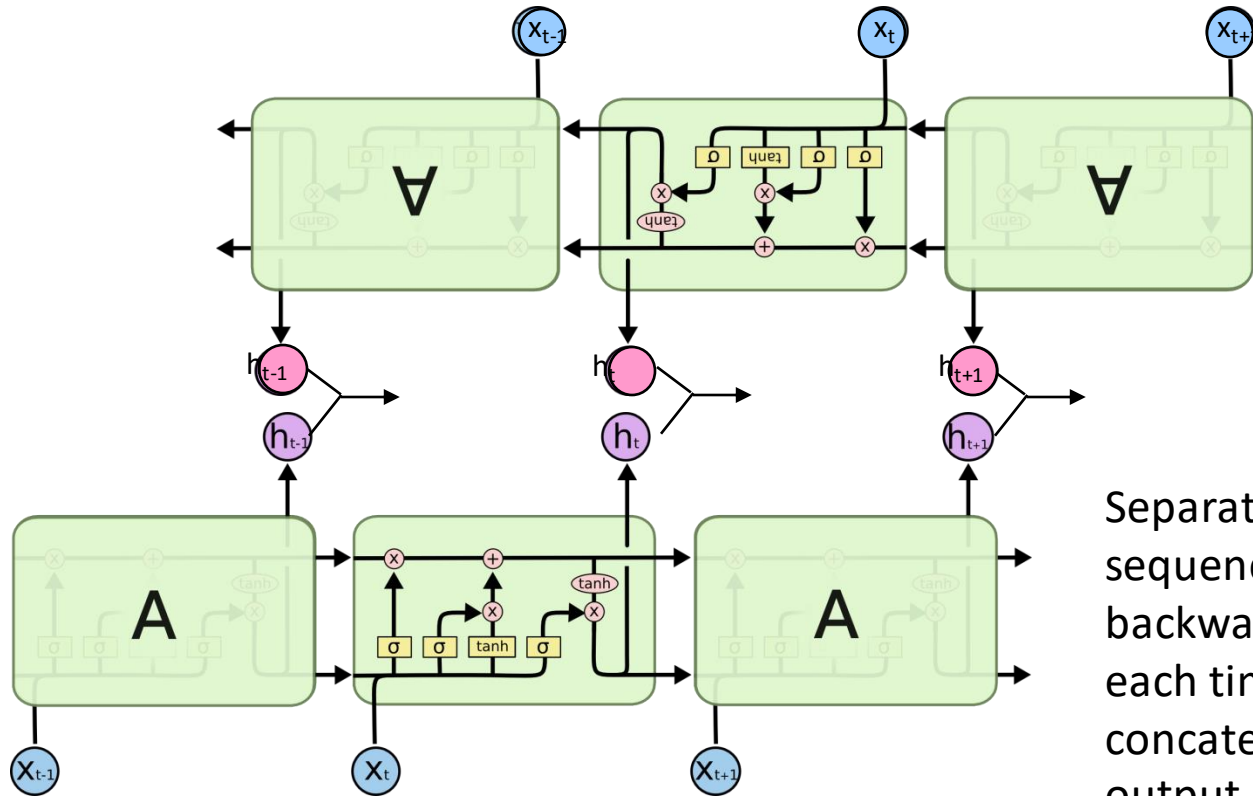


POS Tagging
Language Modeling

Successful Applications of LSTM

- Speech recognition: Language and acoustic modeling
- Sequence labeling
 - POS Tagging
[https://www.aclweb.org/aclwiki/index.php?title=POS_Tagging_\(State_of_the_art\)](https://www.aclweb.org/aclwiki/index.php?title=POS_Tagging_(State_of_the_art))
 - NER
 - Phrase Chunking
- Neural syntactic and semantic parsing
- Image captioning: CNN output vector to sequence
- Sequence to Sequence
 - Machine Translation (Sustkever, Vinyals, & Le, 2014)
 - Video Captioning (input sequence of CNN frame outputs)


Bi-directional LSTM (Bi-LSTM)



Separate LSTMs process sequence forward and backward and hidden layers at each time step are concatenated to form the cell output.

Gated Recurrent Unit (GRU)

- Alternative RNN to LSTM that uses fewer gates ([Cho, et al., 2014](#))
 - Combines forget and input gates into “update” gate.
 - Eliminates cell state vector



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$


$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

GRU vs. LSTM

- GRU has significantly fewer parameters and trains faster.
- Experimental results comparing the two are still inconclusive, many problems they perform the same, but each has problems on which they work better.



Attention

- For many applications, it helps to add “attention” to RNNs.
- Allows network to learn to attend to different parts of the input at different time steps, shifting its attention to focus on different aspects during its processing.
- Used in image captioning to focus on different parts of an image when generating different parts of the output sentence.
- In MT, allows focusing attention on different parts of the source sentence when generating different parts of the translation.

Conclusions

- By adding “gates” to an RNN, we can prevent the vanishing/exploding gradient problem.
- Trained LSTMs/GRUs can retain state information longer and handle long-distance dependencies.
- Recent impressive results on a range of challenging NLP problems.