



UNIVERSITY OF
ARKANSAS

Graph Neural Network

Jiahui Chen

Department of Mathematical Sciences

University of Arkansas

Reference: Kipf's slides

Deep Learning


IMAGENET



Speech data

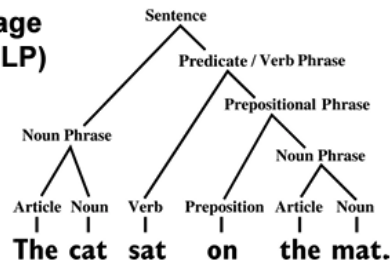


Grid games




Natural language
processing (NLP)

...



Deep neural nets that exploit:


- translation equivariance (weight sharing)
- hierarchical compositionality




Graph Structured Data

A lot of real-world data does not “live” on grids

Social networks
Citation networks
Communication networks
Multi-agent systems




Protein interaction networks




Knowledge graphs



Molecules




Road maps



How to apply the standard deep learning architectures?

Graph Neural Networks (GNNs)

The bigger picture:





Notation: $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

- Adjacency matrix $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix $\mathbf{X} \in \mathbb{R}^{N \times F}$

Main idea: Pass messages between pairs of nodes & agglomerate

Graph Convolutional Network



Update for a single pixel:

- Transform messages individually $\mathbf{W}_i \mathbf{h}_i$
- Add everything up $\sum_i \mathbf{W}_i \mathbf{h}_i$


$\mathbf{h}_i \in \mathbb{R}^F$ are (hidden layer) activations of a pixel/node

Full update:


$$\mathbf{h}_4^{(l+1)} = \sigma \left(\mathbf{W}_0^{(l)} \mathbf{h}_0^{(l)} + \mathbf{W}_1^{(l)} \mathbf{h}_1^{(l)} + \cdots + \mathbf{W}_8^{(l)} \mathbf{h}_8^{(l)} \right)$$

Graph Convolutional Network

Consider this
undirected graph:



Calculate update
for node in red:



**Update
rule:**

$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

Desirable properties:


- Weight sharing over all locations
- Invariance to permutations
- Linear complexity $O(E)$
- Applicable both in transductive and inductive settings

Limitations:

- Requires gating mechanism/residual connections for depth
- Only indirect support for edge features

\mathcal{N}_i : neighbor indices c_{ij} : norm. constant
(fixed/trainable)

GCNs with Edge Embeddings



Formally:

$$v \rightarrow e : \mathbf{h}_{(i,j)}^l = f_e^l([\mathbf{h}_i^l, \mathbf{h}_j^l, \mathbf{x}_{(i,j)}])$$

$$e \rightarrow v : \mathbf{h}_j^{l+1} = f_v^l([\sum_{i \in \mathcal{N}_j} \mathbf{h}_{(i,j)}^l, \mathbf{x}_j])$$


Pros:

- Supports edge features
- More expressive than GNN
- Support sparse matrix

Cons:

- Need to store intermediate edge-based activations
- Difficult to implement with subsampling
- In practice limited to small graph

Graph Neural Network with Attention



[Figure from Veličković et al. (ICLR 2018)]

$$\vec{h}'_i = \sigma \left(\frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$

$$\alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{a}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{a}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_k] \right) \right)}$$

Pros:


- No need to store intermediate edge-based activation vectors (when using dot-product attn.)
- Slower than GCNs but faster than GNNs with edge embeddings

Cons:

- (Most likely) less expressive than GNNs with edge embeddings
- Can be more difficult to optimize


Graph Pooling

The bigger picture:



Graph Pooling:

- Max-pooling
- Min-pooling
- Mean-pooling



Main idea: Pass messages between pairs of nodes & agglomerate